

The Anticipation Content of Fiscal Policy: UK Tax Measures, 1945–2019

Alessandro Conway

Working Paper AC-WP-2021-01

Version 2.0 (revised September 2026)

Not peer reviewed. Comments welcome.

Abstract

Narrative tax shocks are dated by implementation, and the resulting series are pooled across decades. Using 2,252 UK tax measures announced between 1945 and 2019, I show that the interval between announcement and implementation has changed enough that such a series mixes two economically different objects. The share of the tax impulse announced more than 120 days before taking effect rises from 0.18 in 1945–79 to 0.70 since 2000, and the unanticipated component of the quarterly impulse falls roughly fivefold. Computing the same quantity from the century-long dataset of Cloyne et al. (2025), on their coding and their threshold, reproduces the series and extends it back to 0.02 in 1920–44. The interval is not incidental. Within the same Budget, National Insurance changes are 39 percentage points more likely than excise duties to arrive with long notice, the fiscal-year calendar produces long notice only when the Budget falls late in the year, and tax rises are four times more likely than cuts to take effect after an election. The rise dates to the early 1990s, coinciding with a reform of the Budget process justified by scrutiny and predictability and never assessed for macroeconomic consequence. Crises do not visibly shorten it.

Keywords: narrative tax shocks, fiscal foresight, tax policy making, announcement effects, fiscal multipliers

JEL classification: D84, E62, E65, H20, H30

Suggested citation: Conway, Alessandro (2021). “The Anticipation Content of Fiscal Policy: UK Tax Measures, 1945–2019.” Working Paper AC-WP-2021-01, Version 2.0.

© Alessandro Conway 2026. Licensed under Creative Commons Attribution 4.0 International (CC BY 4.0).

The Anticipation Content of Fiscal Policy: UK Tax Measures, 1945–2019

Alessandro Conway*

First draft December 2021

This version September 2026

Abstract

Narrative tax shocks are dated by implementation, and the resulting series are pooled across decades. Using 2,252 UK tax measures announced between 1945 and 2019, I show that the interval between announcement and implementation has changed enough that such a series mixes two economically different objects. The share of the tax impulse announced more than 120 days before taking effect rises from 0.18 in 1945–79 to 0.70 since 2000, and the unanticipated component of the quarterly impulse falls roughly fivefold. Computing the same quantity from the century-long dataset of Cloyne et al. (2025), on their coding and their threshold, reproduces the series and extends it back to 0.02 in 1920–44. The interval is not incidental. Within the same Budget, National Insurance changes are 39 percentage points more likely than excise duties to arrive with long notice, the fiscal-year calendar produces long notice only when the Budget falls late in the year, and tax rises are four times more likely than cuts to take effect after an election. The rise dates to the early 1990s, coinciding with a reform of the Budget process justified by scrutiny and predictability and never assessed for macroeconomic consequence. Crises do not visibly shorten it.

Keywords: narrative tax shocks, fiscal foresight, tax policy making, announcement effects, fiscal multipliers.

JEL classification: E62, E65, H20, H30.

*Working paper. The first draft dates from December 2021 and was not circulated; this version revises the analysis throughout and cites work published since. I thank James Cloyne for making the underlying narrative dataset public, without which this paper could not exist.

1 Introduction

Two dates attach to every tax measure: the day it is announced and the day it takes effect. The narrative tradition in fiscal policy has recorded both for forty years and has never asked what sits between them. Romer and Romer (2010) identify legislated US tax changes from the contemporaneous record, code their motivation, and date them by when they take effect, and Cloyne (2013) builds the corresponding UK account for 1945 to 2009. That record has since been extended in both directions, back into the interwar period by Cloyne et al. (2023) and forward to 2020 by Hürtgen et al. (2024), so that a narrative account of British tax policy now spans roughly a century and has been used at that length by Cloyne et al. (2025).

Throughout that work the implementation date is a dating convention, an input to the construction of a shock series rather than a quantity of interest. The announcement date enters, when it enters at all, as the means of separating measures that were foreseen from measures that were not.

The foresight literature asks a sharper question of the same dates. Does anticipation invalidate the identification? Yang (2005) shows that policy foresight changes the dynamic response to a tax change. Leeper et al. (2013) establish that foresight places the econometrician behind the agent's information set, leaving a non-invertible representation from which the shock cannot be recovered in the usual way. Mertens and Ravn (2011, 2012) construct separate anticipated and unanticipated US tax series in order to estimate their effects apart. Ramey (2011) makes the parallel argument for government spending, that the timing of the news rather than the timing of the outlay is what identifies the shock. Favero and Giavazzi (2012) and Alesina et al. (2015) move closer to the policy process itself, the latter treating fiscal consolidations as multi-year plans with announced and unexpected components rather than as sequences of independent measures. Of the existing work this comes nearest to the treatment taken here.

What these two bodies of work share is the level at which anticipation is handled. In both, foresight is a property of an individual measure, to be recorded, corrected for, or modelled, and the correction is applied measure by measure within a sample that is treated as homogeneous. Neither assembles the anticipation of individual measures into a property of the policy regime, and neither asks how much of it there is in a given decade, whether it falls disproportionately on particular instruments, or whether the quantity has moved over the life of the sample. The assumption that it has not moved is nowhere defended, because it is nowhere stated.

Set against that is a second literature, entirely separate, which does treat the timing of tax policy as its subject and treats it as procedure. The parliamentary record and the Treasury's own accounts of how tax policy is made document four reforms of the Budget process: the unification of the Budget and the Autumn Statement announced in 1992, the consultation framework of 2010 and 2011, and the single fiscal event of 2016. That record is explicit about lead times. It states a target of sixteen months from announcement to effect.

That record is in no conversation with the macroeconomic literature. It justifies longer lead times by reference to parliamentary scrutiny and to the predictability of the tax system, and contains no assessment of what a longer lead time does to the transmission of a fiscal impulse. Answering that was never its purpose.

Between the two sits the gap this paper occupies. The macroeconomic literature has es-

established that foresight matters and has built the tools to handle it, one measure at a time and within a sample assumed to be stationary in this respect. The institutional record has established that the process generating foresight was deliberately reformed, twice, and says nothing about the consequence. What has not been done is to measure the foresight itself as a time series, to ask whether it is systematically distributed, and to connect its movement to the reforms that produced it.

Doing so on 2,252 UK tax measures across 122 Budget events returns four results and two null findings. The share of measures arriving with at least 120 days’ notice runs at roughly 15 per cent from the 1950s through the 1980s and reaches 65 per cent in the 2010s. Weighted by revenue, the share of the tax impulse announced more than 120 days ahead rises from 0.18 in 1945–79 to 0.70 since 2000, and the unanticipated component of the quarterly impulse falls by a factor of about five. That quantity replicates on data I did not build: computing it from the century-long series published by Cloyne et al. (2025), on their coding and their threshold, gives 0.25, 0.42 and 0.70 across the same eras, and extends the finding back to 0.02 in 1920–44.

The interval is also systematic, in three ways that hold within Budget events, so that measures announced by the same Chancellor on the same afternoon are compared with one another. National Insurance changes are 38.5 percentage points more likely to arrive with long notice than excise duties and income tax 12.8 points, while capital gains tax is 11.2 points less likely. Among measures pinned to early April by the fiscal year, 32 per cent arrive with long notice after a spring Budget and 86 per cent after an autumn one, so the same rule applied to the same measure delivers three weeks or five months of warning depending only on when the Budget falls. And among measures announced six to twenty-four months before a general election, 12.5 per cent of tax rises take effect after the country has voted against 3.1 per cent of tax cuts.

Crises do not shorten the interval on any of three definitions of treatment. Every estimate runs the other way.

Section 2 describes the data, what the source can and cannot record, and the tests that establish the trend is not an artefact of joining two sources. Sections 3 to 6 present the four results and Section 7 the crisis null. Section 8 converts the measure-level findings into the anticipation content of the revenue impulse and reports the external replication. Section 9 dates the change and sets it against the institutional record.

2 Data

2.1 Construction

Cloyne (2013) codes UK tax measures from 1945 to 2009 from Financial Statement and Budget Reports, recording an announcement date, an implementation date, a revenue costing and a motivation classification. I extend the record from Budget 2004 to Autumn 2018 from Treasury scorecards, using the same concepts and the same fields. The two are chained at 1 January 2004, Cloyne’s coding before and mine after, which gives 2,252 datable measures across 122 Budget events.

The announcement date is not the moment agents first learn that a measure is coming. It is

the date of the scored, dated commitment, taken from the Budget report or the scorecard, and 90 per cent of measures carry Budget day itself. Since 2010 the tax policy making framework has placed a consultation stage ahead of that date for many measures, so for those the information set moved earlier than the record shows.

The measured lead is therefore a lower bound on foresight, and the bound is loosest in the period where the measured lead is longest, so the rise documented below understates the rise in true foresight. Every quantity reported is the formal lead time of scored policy, which coincides with foresight for the measures that arrive without prior consultation.

The sample is every datable, non-reversal measure regardless of exogeneity. How long the policy process takes is a descriptive question with no endogeneity to purge, and imposing the identifying restriction that a causal estimate would require discards 41 per cent of Cloyne's datable measures and 69 per cent of mine, for no methodological gain. Exogeneity is carried through the data because Section 8 needs it and because the contrast between exogenous and endogenous measures is itself informative.

2.2 The outcome

The primary outcome is an indicator for a gap of at least 120 days between announcement and implementation. The concept comes from Cloyne (2013), who, following Mertens and Ravn (2012), classifies a tax change as anticipated when its implementation lag exceeds 90 days, a convention the literature has used since. The 120-day threshold used here is a variant of it, widened by one month for a reason specific to the British fiscal calendar.

The candidate outcomes are three. A fiscal-year measure, recording whether implementation falls in a later fiscal year than announcement, is partly artefactual: a Budget delivered in March and implementing on 6 April crosses the boundary on three weeks' notice, and the frequency of that case drifts from 0.72 of such crossings in the 1980s to 0.11 in the 2010s, so the measure changes meaning over the period being studied. The raw lag in quarters fails for the opposite reason, with 42 per cent of its mass at zero. A day count avoids both, and it is the outcome used here.


Within a day count, the choice between 90 and 120 days turns on the same fiscal calendar. A Budget delivered in late November or December that implements on the following 1 or 6 April sits about 110 days out, so a 90-day rule records it as anticipated when it is the routine start of the next fiscal year rather than a discretionary delay. Widening to 120 days removes that case. It is a small correction: 94 measures, 4.2 per cent of the sample, fall between 91 and 119 days, and 32 of them are autumn or winter Budgets pinned to the following April. Table 1 reports the two headline results on all four outcomes, including Cloyne's own 90-day rule, and the sign and significance are the same throughout.

All inference is unweighted and clustered on the Budget event. Revenue weighting reduces the Kish effective sample from 2,252 to 251, so it is used for descriptive magnitudes and not for tests.

Table 1: The headline results on four candidate outcomes

Outcome	2010s against the 1940s		Social security within Budget	
	Effect	z	Effect	z
120 days or more (primary)	+0.483	6.84	+0.385	7.48
90 days or more (Cloyne, 2013)	+0.506	7.96	+0.356	6.93
Crosses the fiscal year	+0.634	9.68	+0.543	11.59
Lag in quarters	+2.917	11.66	+1.823	3.90

The first two outcomes are indicators and the third is a count of quarters, so only the first three columns are on a common scale. Standard errors clustered on the Budget event.

**Figure 1:** Announcement-to-implementation gap. The mass below thirty days is the March Budget to 6 April convention, a fiscal-year crossing on three weeks' notice.

2.3 What a scorecard records

A scorecard records commitments and says nothing about delivery. Across 3,092 rows there is not one case in which an expiry date reflects a later decision rather than a sunset announced at the outset, and every reversal row is announced on or before its parent and implements exactly on the parent's stop date. The source cannot see a measure that was announced and later dropped without ceremony.

Every claim below therefore concerns scored, dated commitments. Where a sunset is announced it is recorded, and those are countable: 10.6 per cent of post-2004 measures carry one against 3.6 per cent before.

2.4 Joining two sources

The trend crosses a join at 2004, Cloyne's coding before it and mine after. If the two record an announcement differently, or itemise a Budget differently, the rise is a change in the record and not in behaviour. Six tests say it is not.

Both codings place roughly nine tenths of announcements on the Budget day itself, 0.916 for Cloyne and 0.888 for mine, so the announcement date is the same object in each. The years 2004 to 2009 are coded twice, and on the eleven Budgets both records contain, Cloyne's coding gives a long-notice share of 0.460 and mine 0.496, or 0.627 and 0.615 weighted by revenue, with a paired within-event difference of -0.021 and a t statistic of -0.46 . Cloyne splits those same Budgets 2.8 times more finely, at 32.0 measures per event against 11.5, and still returns the same share. Itemisation practice is not driving the result.

The trend itself holds unweighted, from 0.162 to 0.646, at event level, from 0.195 to 0.618, and revenue-weighted, from 0.285 to 0.740. Holding the 1945–79 instrument mix fixed and letting only the within-instrument rates move changes the 2010s figure by 0.019, so the rise is not a change in which taxes governments use. And taking the twenty largest measures of each decade, an equal count that itemisation cannot touch, gives 0.25, 0.10, 0.20, 0.30, 0.20, then 0.70, 0.80 and 0.80.

The decisive test is that the post-1990 step is $+0.206$ ($z = 5.18$) within Cloyne's coding alone, and $+0.175$ ($z = 4.43$) with instrument controls. No join can produce a break that lies entirely on one side of it.

3 The rise in advance notice

Table 2 reports the share of measures arriving with at least 120 days' notice, by decade of announcement. The series is flat through the 1980s and rises from the 1990s, and no decade cell rests on a single Budget, the largest event accounting for between 11 and 25 per cent of a decade's measures.

4 Instrument determines notice

Comparing measures announced in the same Budget removes the Chancellor, the political situation and the macroeconomic conditions of the day. What remains is an association between

Table 2: Share of measures with at least 120 days between announcement and implementation

	1940s	1950s	1960s	1970s	1980s	1990s	2000s	2010s
Share	0.162	0.039	0.129	0.182	0.145	0.363	0.393	0.646
Measures	80	102	178	203	394	534	377	384
Budget events	6	13	17	21	13	17	22	20




Figure 2: The long-notice share in five-year blocks.

instrument and notice, since measures within a Budget still differ in administrative, legislative and statutory-calendar characteristics, and Section 5 shows that the calendar is doing much of the work. Table 3 reports those effects relative to excise duties. The joint test gives $F = 11.7$ on 10 degrees of freedom, and R^2 rises from 0.274 with Budget fixed effects alone to 0.314 once instrument is added.

Table 3: Instrument effects on the long-notice indicator, Budget-event fixed effects

Instrument	Effect	z	p
Social security (NICs)	+0.385	7.48	0.000
Property, recurrent	+0.148	2.03	0.042
Income	+0.128	3.46	0.001
VAT	+0.096	2.27	0.023
Corporate	+0.036	0.83	0.406
Capital gains	−0.112	−3.25	0.001

Reference category is excise and duties. Oil, property transaction, inheritance and other are estimated and insignificant. Standard errors clustered on the Budget event.

The specification carries 128 parameters against 118 clusters, so the cluster-robust variance matrix is rank deficient and the standard errors are suspect in principle. A pairs cluster bootstrap over 600 replications, resampling whole Budget events, returns standard deviations of 0.046, 0.076, 0.035, 0.043 and 0.034 for social security, recurrent property, income, VAT and capital gains, against cluster-robust standard errors of 0.051, 0.073, 0.037, 0.042 and 0.034. A within-transformed specification, demeaning by Budget event and so carrying ten parameters against the same 118 clusters, returns identical point estimates and standard errors of 0.050, 0.071, 0.036, 0.041 and 0.033. The result is not an artefact of the estimator, with one qualification: recurrent property is marginal on every version of the variance, and 74 of the 600 bootstrap samples contain no recurrent property measure at all, so that row is the least secure of the five.

5 The fiscal-year calendar

Income tax and National Insurance changes are pinned to early April by the fiscal year, so they sit at the top of Table 3. That pinning produces long notice only when the Budget falls late in the year, and the point can be tested directly on the 1,050 measures implementing between 1 and 7 April. Among those measures the long-notice rate is 0.32 after a spring Budget and 0.86 after an autumn one. The raw difference of 0.54 overstates the effect, because autumn Budgets are concentrated in the later decades when notice was longer for every instrument; conditioning on decade gives +0.365 with $z = 5.14$, clustered on the Budget event. The rule is the same, the measure is the same, and the amount of warning differs by five months because of where the Budget sits.

This explains variation across measures and not the trend. Adding calendar controls raises R^2 from 0.155 to 0.214 and absorbs 9 per cent of the 2010s decade effect. The implementation date has also migrated within April, the 6 April share falling from 0.47 in the 1970s to 0.06 in the 2010s while the 1 April share rises from 0.10 to 0.53, which is a change in administrative




Figure 3: Instrument effects with clustered 95 per cent confidence intervals.

convention rather than in notice.

6 Tax rises are scheduled past elections

Among measures announced between six and twenty-four months before a general election, 3.1 per cent of tax cuts and 12.5 per cent of tax rises take effect on the far side of the vote. Within Budget the difference is +0.075 ($z = 2.37$, $p = 0.018$) on 751 measures across 49 events, and on the wider window of nought to twenty-four months it is +0.060 ($z = 2.26$, $p = 0.024$).

The pattern is a descriptive association, and three things stand between it and a reading as deliberate scheduling. Election dates were not fixed before 2011, so the vote a measure is compared against was itself partly a government choice, and the causation can run from the fiscal calendar to the election date rather than the other way. The date was in any case not known when most of these measures were announced, so the test conditions on a realised election rather than an expected one. Budget fixed effects then leave within-Budget differences in instrument mix and statutory calendar untouched, and both correlate with sign, since rises lean towards the fiscal-year-pinned instruments of Section 5 and those land further out by construction.

Restricting to the fixed-term period would answer the first two, and eight years of it will not carry the test. The pattern stands, and its interpretation is open.

7 Crises do not shorten the interval

Crises do not shorten the interval in these data. Taking the years 1974–76, 1992–93 and 2008–09 as crisis windows and conditioning on decade, the estimate is +0.047 on 753 endogenous measures of which 159 are in a window ($p = 0.57$, 95 per cent interval -0.113 to $+0.206$), and

Why instruments differ: the fiscal-year clock

Measures pinned to early April, by Budget season.
The same measure is three weeks from a March Budget and five months from a November one.




Figure 4: April-pinned measures by Budget season.

Tax rises are scheduled past the next election

Share of measures whose implementation date falls after the next general election, for measures announced 6-24 months before that election




Figure 5: Rises and cuts landing after the next general election.

+0.070 on 1,208 exogenous measures of which 84 are in a window ($p = 0.65$, interval -0.233 to $+0.372$). Both are positive, neither is significant, and earlier tests on the population of crisis-response measures and on years containing three or more fiscal events are null or wrong-signed.

The intervals rule out a crisis shortening of more than 11 percentage points on the endogenous sample and 23 on the exogenous one, and the smallest true shortening this design would detect four times in five is 23 and 43 points respectively, so a moderate crisis effect could sit inside those bounds.

The descriptive record runs the other way and runs harder. Of the 22 exogenous measures announced in 2008–09, 95.5 per cent carried long notice against 34.4 per cent of the 195 announced elsewhere in that decade. Twenty-two measures will not carry a test, and the contrast points the same way as the estimates. Crises in this sample bring more pre-commitment, which is the opposite of how a capacity held in reserve for emergencies would behave.

8 The anticipation content of the fiscal impulse

Sections 3 to 6 count measures. The object other researchers use is the revenue impulse, and the two need not move together, since a Budget can pre-announce many small measures and spring one large one. Table 4 reports the revenue-weighted share of the impulse, dated by implementation, that was announced at least 120 days earlier. Measures implementing after 2019 are dropped, since they are the long-lead measures still pending when the sample ends.

Because these quantities are the ones another researcher would want to reconstruct, the aggregation rules are stated here. Write v_i for the costing of measure i , signed so that a rise is positive, and ℓ_i for the indicator that its announcement preceded implementation by at least 120 days. The anticipation share in Table 4 is $\sum_i |v_i| \ell_i / \sum_i |v_i|$, taken over measures whose implementation date falls in the era, and the revenue-weighted lead is $\sum_i |v_i| d_i / \sum_i |v_i|$ with d_i the gap in days censored below at zero. Both use absolute costings, so a rise and a cut of equal size contribute equally and neither nets against the other.

Costings are the peak-year figure at announcement, expressed as a share of that year’s nominal GDP and never revised, so subsequent re-costings do not enter. A measure with several implementation dates is dated by the first. A measure with a pre-announced sunset is counted once, at implementation.

The quarterly decomposition below is a different object and nets within the quarter. Let F_t be the sum of v_i over measures implementing in quarter t with $\ell_i = 1$, and U_t the same sum over those with $\ell_i = 0$, each divided by nominal GDP in that quarter. The paired figures reported by subperiod are $\overline{|F_t|}$ and $\overline{|U_t|}$, the mean absolute net quarterly impulse in each component, and the dispersion figure is the standard deviation of the signed U_t . The two statistics answer different questions, so the shares implied by them differ slightly: the first asks what fraction of the gross flow was pre-announced, the second how large a typical quarter’s net surprise was.

Decomposing the quarterly impulse into a foreseen and an unforeseen component, as a share of nominal GDP, gives 0.041 and 0.077 in 1945–79, 0.049 and 0.083 in the 1980s, 0.051 and 0.030 from 1990 to 2003, and 0.030 and 0.015 from 2004 to 2018. What has collapsed is the surprise. The unforeseen impulse falls by a factor of about five and its standard deviation from 0.240

Table 4: Share of the tax impulse announced at least 120 days ahead

Sample	1945–79	1980–99	2000–19
All measures	0.184	0.384	0.699
Household, exogenous	0.202	0.364	0.815
Revenue-weighted lead, days (all measures)	102	147	322

to 0.033 per cent of GDP, while the foreseen impulse falls by much less, so the foreseen share rises. The collapse occurs between the 1980s and the period from 1990 to 2003, a window lying entirely within Cloyne’s coding, so it cannot be a property of my own data.

8.1 External replication

The result above is computed on a series I assembled, and it can be computed on someone else’s. Cloyne et al. (2025) use a narrative UK dataset covering roughly 1918 to 2020, and their replication package publishes the quarterly exogenous shock series in two variants, a baseline containing all narrative shocks and an unanticipated variant restricted to measures implemented within Cloyne’s 90-day window. Writing B_t for the baseline series and U_t for the unanticipated variant, the anticipation share for an era is $1 - \sum_t |U_t| / \sum_t |B_t|$, summing absolute quarterly values over the quarters in the era, so signs, within-quarter netting and zero denominators do not arise. This is an anticipation share constructed by other researchers, on their own coding, over a longer sample, using their own threshold.

Table 5: Anticipation share of the tax impulse, two independent datasets

Era	Cloyne et al. (2025), exogenous, 90-day rule	This paper, all measures, 120-day rule
1920–44	0.022	not covered
1945–79	0.248	0.184
1980–99	0.420	0.384
2000–20 / 2000–19	0.697	0.699

The two columns are built differently. Their series is exogenous measures only and runs to 2020; the column from this paper is the all-measures row of Table 4 and stops in 2019. The household-exogenous row of Table 4, which is closer to their sample in construction, gives 0.202, 0.364 and 0.815. The comparison establishes that the level and the shape of the series hold under a different coding, a different threshold and a different sample.

The two agree in shape and in level, and most closely where the sample is densest. Their decade series also contains the break, running 0.286 in the 1980s and 0.614 in the 1990s, which is where Section 9 dates it independently from measure-level data. And their longer sample strengthens the result, because in 1920–44 the anticipation share was 0.022, so on a century view the quantity has risen from almost nothing to roughly four fifths.

The stock of announced but not yet implemented tax change moves with it. The median pending stock rises from 0.026 per cent of GDP in 1945–79 to 0.991 per cent in 2004–18, and the mean number of pending measures from 2.0 to 26.4. The median is the statistic to report, because the mean is dominated by the introduction of VAT, announced in March 1971 and in force in April 1973, which alone placed roughly 11 per cent of GDP in the pipeline. That

measure is the point in miniature. Long pre-announcement was once exceptional enough to be a landmark, and it is now routine.

The consequence for applied work follows. A multiplier estimated on a UK narrative series pooled across the post-war period averages a regime in which tax changes mostly arrived as surprises with one in which they mostly did not, and the mix changes monotonically over the sample.

Whether that pooling biases the estimate, and in which direction, depends on how news enters the estimator and on whether anticipated and unanticipated UK measures generate different macroeconomic responses, which is the companion paper’s question. What this paper settles is that the information content of the series is not stationary, and that is a property of the instrument. The announcement-dated and implementation-dated series built here are supplied as separate instruments so the two can be estimated apart, and Leeper et al. (2013) give the reasons for expecting the distinction to matter.

9 Dating the change

A grid search over single break years maximises fit at 1993, with an R^2 of 0.142, and the annual series is sharp around it, running 0.047 in 1992, 0.372 in 1993, 0.434 in 1994 and 0.600 in 1995. The break year is searched rather than assumed, so its nominal fit is not a test statistic and no p -value is attached to it; the five best-fitting years span 1989 to 1995, so the change dates to the early 1990s rather than to a single year. Table 6 then tests each named candidate locally, on the eight years either side, taking the date as given.

Table 6: Institutional candidates, local windows of eight years either side

Candidate	Year	Estimate	z	p
Unified autumn Budget, first held Nov. 1993	1993	+0.288	6.96	< 0.0001
Return to a spring Budget	1997	+0.026	0.36	0.716
Tax policy making framework	2010	+0.257	3.20	0.001
Draft-clause consultation	2011	+0.157	1.80	0.072
Single fiscal event	2017	+0.108	1.20	0.231

Two of these hold under revenue weighting and two do not, which matters because the object of Section 8 is the money rather than the count of measures. Repeating the exercise on the annual revenue-weighted long-notice share, the grid search again picks 1993, with a higher R^2 of 0.431 and the same 1989 to 1995 spread, and the local test at 1993 strengthens to +0.429 ($t = 4.02$, $p = 0.001$). The 2010 candidate does not hold: +0.190 with $p = 0.116$. The 2010 framework moved the number of measures given long notice without detectably moving the share of the revenue, so it belongs in the account of the trend but not in the account of the impulse. Only the 1993 date is supported on both.

The 1993 step is not a calendar artefact. Restricting to spring Budgets alone, the long-notice share still runs 0.143 in the 1980s, 0.275 in the 1990s and 0.657 in the 2010s, with a post-1990 step of +0.275 ($z = 5.56$), and the autumn-minus-spring gap runs +0.21, +0.17 and -0.03 across the last three decades, so the calendar channel closes while the level continues to rise.

The rejection of 1997 makes the same point from the other side. The Budget returned to spring and the long-notice share did not return with it.

Both dates correspond to reforms of the Budget process. Norman Lamont announced the first in the Budget Statement of 10 March 1992, that “next year’s Budget will be the last spring Budget” and that “from then on the annual Budget will be in December, and it will cover not just taxation but also public expenditure”, adding that the 1994 Finance Bill would be presented in January rather than April, which moved the legislative timetable as well as the announcement. The date was brought forward to November in the Budget Statement of 16 March 1993, and the first unified Budget was delivered by Kenneth Clarke on 30 November 1993.

Budget 2010 set out predictability, stability and simplicity as objectives and published proposals “to improve the way tax policy is made to support these objectives”, and the consultation response of December 2010 recorded that confirming most intended tax changes “at least three months ahead of publication of the Bill supports predictability in the tax system”. The move to a single fiscal event in 2016 was made “to promote certainty and simplicity within the tax system”, and the Treasury’s current statement of the resulting timetable is that “most policies will be announced at least 16 months before they come into effect at the start of the next tax year”.

The stated objectives across 1992, 2010, 2016 and 2017 are parliamentary and external scrutiny, predictability, stability, certainty, simplicity, and time for stakeholders to comment. None is macroeconomic. Each reform was argued for in public and chosen deliberately, and in none of those arguments does the anticipation content of fiscal policy appear.

The design cannot establish that the reforms caused the shift: there is one country, one series, a searched break date, candidate reforms in adjacent years that no test on annual data can separate, and a strong secular trend running underneath. What it establishes is that the change coincides with the 1993 Budget-process reform, that the coincidence holds under revenue weighting and the removal of the calendar channel, and that the 1997 reversal of the Budget season did not reverse it. Whatever caused the change, the anticipation content of British fiscal policy is now different from what the literature’s pooled samples assume, and it moved without anyone intending it to.

10 Conclusion

The interval between announcing a British tax change and its taking effect has roughly quadrupled since the middle of the last century, the unanticipated component of the tax impulse has fallen by about five to one, and on a century view the anticipation share has risen from close to zero to roughly four fifths. The interval depends on which instrument is chosen, on where the Budget falls in the calendar, and on whether an election is approaching, and it does not respond to crises.

The distinction between anticipated and unanticipated tax measures, and the tools for handling foresight, come from the literature. What this paper adds is the treatment of anticipation as a property of the policy regime, measured as a time series, shown to be systematically dis-

tributed, and dated to identifiable reforms. The ingredients sat in a public replication archive, where they served as a robustness check.

Two consequences follow. For applied work, a narrative tax series pooled across the post-war period is not a series of comparable objects, and the announcement-dated and implementation-dated series built here let a researcher estimate the two apart. For policy, the timing of a tax change carries economic content and is at present a residual of instrument choice and Budget scheduling, with a political tilt at the margin.

Notice is worth having only to a household that can act on it, and acting on it takes savings, flexible income or paid advice. A large increase in notice is therefore a large increase in a benefit that is unevenly distributed by construction. Whether the unevenness matters for household behaviour cannot be tested in this dataset, and the simplest version of the claim fails on what can be tested: the gap in notice between regressive and progressive instruments was significant between 1980 and 1999 and is not significant today. The raw difference persists, with 43 per cent of regressive measures arriving on under a month's warning against 33 per cent of progressive ones. The companion paper takes up who receives the notice.

Britain reformed the timing of its tax policy for reasons of parliamentary scrutiny and legal certainty, and in doing so transformed the anticipation content of fiscal policy. That consequence appears in none of the documents that made the case for the reforms.

References

- Alesina, A., Favero, C. and Giavazzi, F. (2015), 'The output effect of fiscal consolidation plans', *Journal of International Economics* **96**, S19–S42.
- Cloyne, J. (2013), 'Discretionary tax changes and the macroeconomy: new narrative evidence from the United Kingdom', *American Economic Review* **103**(4), 1507–1528.
- Cloyne, J., Dimsdale, N. and Postel-Vinay, N. (2023), 'Taxes and growth: new narrative evidence from interwar Britain', *Review of Economic Studies* .
- Cloyne, J., Hürtgen, P. and Dimsdale, N. (2025), 'Are tax cuts contractionary at the zero lower bound? Evidence from a century of data', *Journal of Political Economy* **133**(2), 568–603.
- Favero, C. and Giavazzi, F. (2012), 'Measuring tax multipliers: the narrative method in fiscal VARs', *American Economic Journal: Economic Policy* **4**(2), 69–94.
- Hürtgen, P., Cloyne, J., Dimsdale, N. and Postel-Vinay, N. (2024), Tax changes in the United Kingdom 2009–2020: a new narrative account and dataset, Technical report, SSRN Working Paper 4909051.
- Leeper, E. M., Walker, T. B. and Yang, S.-C. S. (2013), 'Fiscal foresight and information flows', *Econometrica* **81**(3), 1115–1145.
- Mertens, K. and Ravn, M. O. (2011), 'Understanding the aggregate effects of anticipated and unanticipated tax policy shocks', *Review of Economic Dynamics* **14**(1), 27–54.

- Mertens, K. and Ravn, M. O. (2012), ‘Empirical evidence on the aggregate effects of anticipated and unanticipated US tax policy shocks’, *American Economic Journal: Economic Policy* **4**(2), 145–181.
- Ramey, V. A. (2011), ‘Identifying government spending shocks: it’s all in the timing’, *Quarterly Journal of Economics* **126**(1), 1–50.
- Romer, C. D. and Romer, D. H. (2010), ‘The macroeconomic effects of tax changes: estimates based on a new measure of fiscal shocks’, *American Economic Review* **100**(3), 763–801.
- Yang, S.-C. S. (2005), ‘Quantifying tax effects under policy foresight’, *Journal of Monetary Economics* **52**(8), 1557–1568.